

Project 4

Yao Liu

Concurrent nodes join

- In Project 3, nodes join one by one.
 - A join() RPC call won't return until all fingertable updates, file transfers have taken place.
 - If implemented correctly, all nodes will have the correct fingertable entries.
- In Project 4, nodes may join the DHT concurrently.
 - We cannot guarantee nodes' fingertables are all correct at all times.
 - but all fingertable entries will be updated over time and eventually set to their correct values

join()

- The join() function will be different from Project 3.
- It does not try to create a fingertable with all entries set.
 - fingertable entries (except the first one) are updated by the fixFinger() function.
- Instead, it only uses an existing DHT node to find its successor, i.e., the first entry in the fingertable
- notify() its successor to update the predecessor

getNodeSucc() and getNodePred()

- `getNodeSucc()` returns the node's successor, which is the first entry in the node's fingertable.
- `getNodePred()` returns the node's predecessor
 - Initially, the node does not know about its predecessor, and this may return NULL (or a structure indicating the predecessor has not been set).
- `getNodeSucc()` and `getNodePred()` may not return the correct value, but the successor and predecessor are periodically updated by the `stabilize()` and `notify()` functions.

stabilize()

- A node n periodically checks if $\text{pred}(\text{succ}(n))$ belongs to the interval: $(n, \text{succ}(n))$
 - $\text{succ}(n)$ is the first entry in n 's fingertable
 - $\text{pred}(m)$ returns node m 's predecessor by calling $\text{getNodePred}()$ on node m .
- If yes, there exists a node p between n and $\text{succ}(n)$
 - Set p to its own successor (set p to the first entry in its fingertable)
 - Use the **notify()** function to notify p that it (node n) is its predecessor
- If the predecessor of $\text{succ}(n)$ is not set
 - **notify()** $\text{succ}(n)$ that node n is its predecessor

notify()

- `notify()` takes one node ID, q , as input argument
- Node n that receives the `notify()` RPC call
 - will set q as its predecessor if its predecessor is currently not set.
 - will set q as its predecessor if q is in the interval of (n 's current predecessor, n)

- `stabilize()` and `notify()` are called periodically and eventually, all predecessors and successors of all nodes in the DHT are eventually set to their correct values


Use `fixFingers()` to update fingertable entries

- `stabilize()` and `notify()` correct nodes' predecessors and successors
- Nodes' fingertable entries are corrected by `fixFingers()`
- A node periodically selects a fingertable entry at random and updates it using `findSucc(id)`.

Test your implementation


- You will need to write a test program to check if after a sufficient amount of time
 - all nodes' successors are correct
 - all nodes' predecessors are correct
 - the number of correct fingertable entries increases
 - the number of incorrect fingertable entries set by `fixFingers` first increases then decreases (unset fingertable entries are not counted as incorrect)
- Use `cmp_fingertables` to determine if fingertable entries are correct

The number of correct successors



64 DHT nodes in total



The number of correct/incorrect fingertable entries



64 DHT nodes in total

Extra 30% credit

- Extend your test program to gauge the performance of the stable Chord protocol.
- Need to output two additional pieces of information:
 - the average number of correct successor nodes returned by findSucc().
 - the average number of calls to findPred() for each correct successor returned
 - Use the optional count field in the NodeID structure.



64 DHT nodes in total